

Machine Learning 2 — Homework 1

Maurice Frank

11650656

maurice.frank@posteo.de

September 9, 2019

Problem 1.

We have: $x \in \mathbb{R}^n$, $p(x) = \mathcal{N}(x|\mu_x, \Sigma_x)$ and $z \in \mathbb{R}^n$, $p(z) = \mathcal{N}(z|\mu_z, \Sigma_z)$. And:

$$\begin{aligned} \mathbf{y} &= \mathbf{x} + \mathbf{z} \\ \implies \\ \mathbb{E}[\mathbf{y}] &= \mathbb{E}[\mathbf{x} + \mathbf{z}] \\ &= \mathbb{E}[\mathbf{x}] + \mathbb{E}[\mathbf{z}] \\ &= \mu_x + \mu_z \\ \text{cov}(\mathbf{y}) &= \mathbb{E}[(\mathbf{x} + \mathbf{y})^2] - (\mathbb{E}[\mathbf{x} + \mathbf{y}])^2 \\ &= \mathbb{E}[\mathbf{x}^2] + 2\mathbb{E}[\mathbf{x}\mathbf{y}] + \mathbb{E}[\mathbf{y}^2] - \mathbb{E}[\mathbf{x}]^2 - 2\mathbb{E}[\mathbf{x}]\mathbb{E}[\mathbf{y}] - \mathbb{E}[\mathbf{y}]^2 \\ &= \mathbb{E}[\mathbf{x}^2] - \mathbb{E}[\mathbf{x}]^2 + \mathbb{E}[\mathbf{y}^2] - \mathbb{E}[\mathbf{y}]^2 \\ &= \text{cov}(\mathbf{x}) + \text{cov}(\mathbf{y}) \end{aligned}$$

Problem 2.

We have: $x \in \mathbb{R}^D$, $p(x) = \mathcal{N}(x|\mu, \Sigma)$ and observe $\chi = \{x_1, \dots, x_N\}$. Further it is $x_i \sim \mathcal{N}(\mu, \Sigma)$ and $\mu \sim \mathcal{N}(\mu_0, \Sigma_0)$.

1.

The likelihood is:

$$\begin{aligned}
p(\chi|\mu, \Sigma) &= \prod_{i=0}^N p(x_i|\mu, \Sigma) \\
&= \prod_{i=0}^N \mathcal{N}(x_i|\mu, \Sigma) \\
&= \prod_{i=0}^N (2\pi)^{-\frac{D}{2}} |\Sigma|^{-\frac{1}{2}} \exp\left(-\frac{1}{2}(x_i - \mu)^T \Sigma^{-1}(x_i - \mu)\right)
\end{aligned}$$

2.

Using Bayes we get the posterior:

$$\begin{aligned}
p(\mu|\chi, \Sigma, \mu_0, \Sigma_0) &= \frac{p(\chi|\mu, \Sigma) \cdot p(\mu|\mu_0, \Sigma_0)}{p(\chi)} \\
&\sim \prod_{i=0}^N \mathcal{N}(x_i|\mu, \Sigma) \cdot \mathcal{N}(\mu_0, \Sigma_0)
\end{aligned}$$

3.

$$\begin{aligned}
p(\mu|\chi, \Sigma, \mu_0, \Sigma_0) &\sim \prod_{i=0}^N \mathcal{N}(x_i|\mu, \Sigma) \cdot \mathcal{N}(\mu_0, \Sigma_0) \\
&\sim \exp \left[-\frac{1}{2}(\mu - \mu_0)^T \Sigma_0^{-1}(\mu - \mu_0) - \frac{1}{2} \sum_{i=0}^N (x_i - \mu)^T \Sigma^{-1}(x_i - \mu) \right] \\
&\sim \exp \frac{1}{2} \left[-\mu^T (\Sigma_0^{-1} + N \cdot \Sigma^{-1}) \mu + 2\mu^T (\Sigma_0^{-1} \mu_0 + \Sigma^{-1} \sum_{i=0}^N x_i) - \mu_0^T \Sigma_0^{-1} \mu_0 - \sum_{i=0}^N x_i^T \Sigma^{-1} x_i \right] \\
&\stackrel{\text{def}}{=} \exp \left[-\mu^T \Sigma_N^{-1} \mu + 2\mu^T \Sigma_N^{-1} \mu_N - \mu_N^T \Sigma_N^{-1} \mu_N \right] \\
&\implies \\
&\Sigma_N^{-1} = \Sigma_0^{-1} + N \cdot \Sigma^{-1} \\
&\quad \wedge \\
2\mu^T (\Sigma_0^{-1} \mu_0 + \Sigma^{-1} \sum_{i=0}^N x_i) &= 2\mu^T \Sigma_N^{-1} \mu_N \\
\mu_N &= \Sigma_N (\Sigma_0^{-1} \mu_0 + \Sigma^{-1} \sum_{i=0}^N x_i)
\end{aligned}$$

With both parameters calculated we have the posterior:

$$\begin{aligned} p(\boldsymbol{\mu}|\boldsymbol{\chi}, \boldsymbol{\Sigma}, \boldsymbol{\mu}_0, \boldsymbol{\Sigma}_0) &= \mathcal{N}(\boldsymbol{\mu}_N, \boldsymbol{\Sigma}_N) \\ &= \mathcal{N}(\boldsymbol{\Sigma}_N(\boldsymbol{\Sigma}_0^{-1}\boldsymbol{\mu}_0 + \boldsymbol{\Sigma}^{-1} \sum_{i=0}^N \boldsymbol{x}_i), \frac{1}{\boldsymbol{\Sigma}_0^{-1} + N \cdot \boldsymbol{\Sigma}^{-1}}) \end{aligned}$$

4.

We are using the log posterior for the MAP estimate:

$$\begin{aligned} \frac{\partial}{\partial \boldsymbol{\mu}} \ln p(\boldsymbol{\mu}|\boldsymbol{\chi}, \boldsymbol{\Sigma}, \boldsymbol{\mu}_0, \boldsymbol{\Sigma}_0) &= 0 \\ &= \frac{\partial}{\partial \boldsymbol{\mu}} - \frac{ND}{2} \ln(2\pi) - \frac{N}{2} \ln |\boldsymbol{\Sigma}| - \frac{1}{2} \sum_{i=0}^N (\boldsymbol{x}_i - \boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1} (\boldsymbol{x}_i - \boldsymbol{\mu}) \\ &\quad - \frac{D}{2} \ln(2\pi) - \frac{1}{2} \ln |\boldsymbol{\Sigma}| - \frac{1}{2} (\boldsymbol{\mu} - \boldsymbol{\mu}_0)^T \boldsymbol{\Sigma}_0^{-1} (\boldsymbol{\mu} - \boldsymbol{\mu}_0) \\ &= - \sum_{i=0}^N \boldsymbol{\Sigma}^{-1} (\boldsymbol{x}_i - \boldsymbol{\mu}) - \boldsymbol{\Sigma}_0^{-1} (\boldsymbol{\mu} - \boldsymbol{\mu}_0) \\ &= - \boldsymbol{\Sigma}^{-1} \sum_{i=0}^N \boldsymbol{x}_i + N \cdot \boldsymbol{\Sigma}^{-1} \boldsymbol{\mu} - \boldsymbol{\Sigma}_0^{-1} \boldsymbol{\mu} + \boldsymbol{\Sigma}_0^{-1} \boldsymbol{\mu}_0 \\ &\implies \\ \boldsymbol{\mu}_{MAP} &= \frac{-\boldsymbol{\Sigma}_0^{-1} \boldsymbol{\mu}_0 + \boldsymbol{\Sigma}^{-1} \sum_{i=0}^N \boldsymbol{x}_i}{N \cdot \boldsymbol{\Sigma}^{-1} - \boldsymbol{\Sigma}_0^{-1}} \end{aligned}$$

Problem 3.

We have a coin with head-proability of μ .

1.

For this Bernoulli experiment the MLE is:

$$\begin{aligned} \mu_{ML} &= \frac{\sum_{i=1}^N x_i}{N} \\ &= 1 \end{aligned}$$

so the probability is 100 per cent.

2.

We assume $\mu \sim \text{Beta}(\mu|, a, b)$:

$$\begin{aligned}
p(\mu|\chi) &= p(\chi|\mu)p(\mu|a, b) \\
&= \prod_{i=0}^N \mu^{x_i} (1-\mu)^{(1-x_i)} \cdot \frac{\mu^{a-1} (1-\mu)^{b-1}}{B(a, b)} \\
&\implies \\
\ln p(\mu|\chi) &= \sum_{i=0}^N x_i \ln \mu + (1-x_i) \ln (1-\mu) + (a-1) \ln \mu + (b-1) \ln (1-\mu) + C \\
&\implies \\
\frac{\partial \ln p(\mu|\chi)}{\partial \mu} &= \sum_{i=0}^N \frac{x_i}{\mu} - \frac{1-x_i}{1-\mu} + \frac{a-1}{\mu} - \frac{b-1}{1-\mu} \\
&\implies \\
0 &= \frac{\sum x_i + a - 1}{\mu_{MAP}} - \frac{\sum (1-x_i) + b - 1}{1 - \mu_{MAP}} \\
\frac{1}{\mu_{MAP}} &= \frac{\sum (1-x_i) + b - 1}{\sum x_i + a - 1} + 1 \\
\mu_{MAP} &= \frac{\sum x_i + a - 1}{N + a + b - 2}
\end{aligned}$$

for our case that is:

$$\mu_{MAP} = \frac{a+2}{a+b+1}$$

3.

Our posterior distribution is also a Beta distribution:

$$\begin{aligned}
p(\mu|\chi, a, b) &= \frac{p(\chi|\mu) \cdot p(\mu|a, b)}{\int_0^1 p(\chi|\mu) \cdot p(\mu|a, b) d\mu} \\
&= \frac{\binom{m+l}{m} \mu^m (1-\mu)^l \cdot \frac{1}{B(a,b)} \mu^{a-1} (1-\mu)^{b-1}}{\int_0^1 \binom{m+l}{m} \mu^m (1-\mu)^l \cdot \frac{1}{B(a,b)} \mu^{a-1} (1-\mu)^{b-1} d\mu} \\
&= \frac{\mu^m (1-\mu)^l \cdot \mu^{a-1} (1-\mu)^{b-1}}{\int_0^1 \mu^m (1-\mu)^l \cdot \mu^{a-1} (1-\mu)^{b-1} d\mu} \\
&= \frac{\mu^{m+a-1} (1-\mu)^{l+b-1}}{\int_0^1 \mu^{m+a-1} (1-\mu)^{l+b-1} d\mu} \\
&= \frac{\mu^{m+a-1} (1-\mu)^{l+b-1}}{B(m+a, l+b)} \\
&= \text{Beta}(m+a, b+l)
\end{aligned}$$

Our posterior distribution is a Beta distribution with the prev derived mean:

$$\begin{aligned}
p(\mu|m, l, a, b) &= \text{Beta}(m+a, l+b) \\
&\implies \\
\mathbb{E}[p(\mu|m, l, a, b)] &= \frac{m+a}{m+l+a+b}
\end{aligned}$$

For prior mean and ML estimate we have:

$$\begin{aligned}
\mathbb{E}[\text{Beta}(\mu | a, b)] &= \frac{a}{a+b} \\
\mu_{ML} &= \frac{m}{m+l}
\end{aligned}$$

We can form an linear combination to give the posterior mean:

$$\begin{aligned}
\mathbb{E}[p(\mu|m, l, a, b)] &= \frac{a}{m+l+a+b} + \frac{m}{m+l+a+b} \\
&= \frac{a+b}{a+b} \frac{a}{m+l+a+b} + \frac{m+l}{m+l} \frac{m}{m+l+a+b} \\
&= \mathbb{E}[\text{Beta}(\mu | a, b)] \frac{a+b}{m+l+a+b} + \mu_{ML} \frac{m+l}{m+l+a+b}
\end{aligned}$$

Given that $a+b > 0$ and $m+l > 0$ and thus both coefficient in $(0, 1)$ this shows that the posterior mean is between those two estimates.

Problem 4.

Distributions in the exponential family have the form:

$$p(x|\eta) = h(x) \exp(\phi(\eta)T(x) - A(\eta))$$

1.

(i)

$$\begin{aligned}
 \text{Pois}(k|\lambda) &= \frac{\lambda^k \exp -\lambda}{k!} \\
 &= \frac{1}{k!} \cdot \exp \log \lambda^k \cdot \exp -\lambda \\
 &= \frac{1}{k!} \cdot \exp (k \log \lambda - \lambda) \\
 &\Rightarrow \\
 h(x) &= \frac{1}{k!} \\
 \phi(\eta) &= \log \lambda \\
 T(x) &= k \quad \text{(sufficient statistic)} \\
 A(\eta) &= \lambda
 \end{aligned}$$

(ii)

$$\begin{aligned}
 \text{Gam}(\tau|a, b) &= \frac{1}{\Gamma(a)} b^a \tau^{a-1} e^{-b\tau} \\
 &= \exp (a \log b - \log \Gamma(a)) \cdot \exp ((a-1) \log \tau) \cdot \exp (-b\tau) \\
 &= \exp [(a-1) \log \tau - b\tau - (\log \Gamma(a) - a \log b)] \\
 &\Rightarrow \\
 h(x) &= 1 \\
 \phi(\eta) &= \begin{bmatrix} a-1 \\ -b \end{bmatrix} \\
 T(x) &= \begin{bmatrix} \log \tau \\ \tau \end{bmatrix} \quad \text{(sufficient statistic)} \\
 A(\eta) &= \log \Gamma(a) - a \log(b)
 \end{aligned}$$

(iii)

$$\text{Cauchy}(x|\gamma, \mu) = \frac{1}{\pi\gamma} \frac{1}{1 + \left(\frac{x-\mu}{\gamma}\right)^2}$$

Because of the $1 +$ in the denominator the Cauchy distribution can not be factorized into a canonical form of the exponential family. Thus it is not part of it.

(iv)

$$\begin{aligned}
\text{vonMises}(x|\kappa, \mu) &= \frac{1}{2\pi I_0(\kappa)} \exp(\kappa \cos(x - \mu)) \\
&= \exp(\cos(x - \mu)\kappa) \cdot \exp(-\log(2\pi I_0(\kappa))) \\
&= \exp(\kappa(\cos \mu \cos x + \sin \mu \sin x) - \log(2\pi I_0(\kappa))) \\
&\implies \\
h(x) &= 1 \\
\phi(\eta) &= \begin{bmatrix} \kappa \cos \mu \\ \kappa \sin \mu \end{bmatrix} \\
T(x) &= \begin{bmatrix} \cos x \\ \sin x \end{bmatrix} && \text{(sufficient statistic)} \\
A(\eta) &= \log(2\pi I_0(\kappa))
\end{aligned}$$

2.

(i)

$$\begin{aligned}
\mathbb{E}[\text{Pois}(k|\lambda)] &= \sum_{k=0}^{\infty} k e^{-\lambda} \frac{\lambda^k}{k!} \\
&= e^{-\lambda} \sum_{k=0}^{\infty} k \cdot \frac{\lambda^k}{k!} \\
&= e^{-\lambda} \sum_{k=0}^{\infty} k \cdot \frac{\lambda \cdot \lambda^{(k-1)}}{k \cdot (k-1)!} \\
&= \lambda e^{-\lambda} \sum_{k=0}^{\infty} \frac{\lambda^{(k-1)}}{(k-1)!} \\
&= \lambda e^{-\lambda} \exp \lambda \\
&= \lambda
\end{aligned}$$

$$\begin{aligned}
\mathbb{E}[(\text{Pois}(k|\lambda) - \lambda)^2] &= e^{-\lambda} \sum_{k=0}^{\infty} (k - \lambda)^2 \frac{\lambda^k}{k!} \\
&= e^{-\lambda} \sum_{k=0}^{\infty} (k^2 - 2k\lambda + \lambda^2) \frac{\lambda^k}{k!} \\
&= e^{-\lambda} ((\lambda^2 + \lambda)e^\lambda - 2\lambda^2 e^\lambda + \lambda^2 e^\lambda) \\
&= \lambda
\end{aligned}$$

(ii)

$$\begin{aligned}
\mathbb{E}[\text{Gam}(\tau|a, b)] &= \frac{b^a}{\Gamma(a)} \cdot \int_0^\infty \tau \cdot \tau^{a-1} e^{-b\tau} d\tau \\
&= \frac{b^a}{\Gamma(a)} \cdot \int_0^\infty \tau^a e^{-b\tau} d\tau \\
&= \frac{b^a}{\Gamma(a)} \cdot \int_0^\infty \left(\frac{\tau}{b}\right)^a e^{-\tau} \frac{1}{b} d\frac{\tau}{b} \\
&= \frac{b^a}{\Gamma(a)} \cdot \frac{1}{b^{a+1}} \int_0^\infty \tau^a e^{-\tau} d\frac{\tau}{b} \\
&= \frac{1}{\Gamma(a)} \cdot \frac{1}{b} \cdot \Gamma(a+1) \\
&= \frac{1}{\Gamma(a)} \cdot \frac{1}{b} \cdot a \cdot \Gamma(a) \\
&= \frac{a}{b}
\end{aligned}$$

$$\begin{aligned}
\text{Var}(\text{Gam}(\tau|a, b)) &= \mathbb{E}[(\text{Gam}(\tau|a, b) - \frac{a}{b})^2] \\
&= \frac{b^a}{\Gamma(a)} \cdot \int_0^\infty \left(\tau - \frac{a}{b}\right)^2 \cdot \tau^{a-1} e^{-b\tau} d\tau \\
&= \frac{b^a}{\Gamma(a)} \cdot \left[\int_0^\infty \tau^{a+2-1} e^{-\tau} \frac{1}{b^{a+2}} d\frac{\tau}{b} - 2\frac{a}{b} \int_0^\infty \tau^{a+1-1} e^{-\tau} \frac{1}{b^{a+1}} d\frac{\tau}{b} \right. \\
&\quad \left. + \left(\frac{a}{b}\right)^2 \int_0^\infty \tau^{a-1} e^{-\tau} \frac{1}{b^a} d\frac{\tau}{b} \right] \\
&= \frac{b^a}{\Gamma(a)} \cdot \left[\frac{\Gamma(a+2)}{b^{a+2}} - 2\frac{a \cdot \Gamma(a+1)}{b \cdot b^{a+1}} + \frac{a^2 \cdot \Gamma(a)}{b^2 \cdot b^a} \right] \\
&= \frac{b^a}{\Gamma(a)} \cdot \left[\frac{(a^2 + a)\Gamma(a)}{b^{a+2}} - 2\frac{a^2 \cdot \Gamma(a)}{b^{a+2}} + \frac{a^2 \cdot \Gamma(a)}{b^{a+2}} \right] \\
&= \frac{a}{b^2}
\end{aligned}$$

3.

Every member of the exponential family has a conjugate prior as such also the Poisson distribution. The conjugate prior to a Poisson is a Gamma which we

can show by computing the posterior:

$$\begin{aligned}
p(\lambda|k, a, b) &= \frac{p(k|\lambda) \cdot p(\lambda|a, b)}{\int p(k)} \\
&= \frac{\frac{\lambda^k e^{-\lambda}}{k!} \frac{1}{\Gamma(a)} b^a \lambda^{a-1} e^{-b\lambda}}{\int \frac{\lambda^{*k} e^{-\lambda^*}}{k!} \frac{1}{\Gamma(a)} b^a \lambda^{*a-1} e^{-b\lambda^*} d\lambda^*} \\
&= \frac{\lambda^{a+k-1} e^{-(b+1)\lambda}}{\int \lambda^{*a+k-1} e^{-(b+1)\lambda^*} d\lambda^*} \\
&= \frac{\lambda^{a+k-1} e^{-(b+1)\lambda}}{\int \left(\frac{\lambda^*}{b+1}\right)^{a+k-1} e^{-\lambda^*} \frac{1}{b+1} d\frac{\lambda^*}{b+1}} \\
&= \frac{\lambda^{a+k-1} e^{-(b+1)\lambda}}{\frac{1}{(b+1)^{a+k}} \int \lambda^{*a+k-1} e^{-\lambda^*} d\frac{\lambda^*}{b+1}} \\
&= \frac{(b+1)^{a+k}}{\Gamma(a+k)} \cdot \lambda^{a+k-1} e^{-(b+1)\lambda} \\
&= \text{Gamma}(a+k, b+1)
\end{aligned}$$